Hedging the AI Singularity: Asset Pricing with Transformative Technology Risk

Andrew Y. Chen

Federal Reserve Board

April 2025*

Abstract

We develop a consumption-based asset pricing model where AI capabilities advance during economic disasters. AI stocks provide hedging value against singularity risk, commanding high valuations. However, when AI's disaster-contingent growth becomes sufficiently extreme, standard asset pricing breaks down, formally capturing concerns about transformative AI destabilizing economic equilibria.

Keywords: Artificial Intelligence, Disaster Risk, Asset Pricing, Technological Singularity

*email:andrew.y.chen@frb.gov. ChatGPT-o1 and Claude-3.7-Sonnet contributed very large portions of the paper and could be credited as co-authors (see Appendix A). I thank Andrei Goncalves for helpful comments. The views expressed herein are those of the authors and do not necessarily reflect the position of the Board of Governors of the Federal Reserve or the Federal Reserve System.

1 Introduction

Recent advances in artificial intelligence have rekindled discussions about transformative technological change. OpenAI's o3 model achieved breakthrough performance on the ARC-AGI benchmark (**chollet2024arc**), while reasoning models from DeepSeek, Anthropic, and Google demonstrate increasingly sophisticated problem-solving capabilities (**deepseek2025deepseek**; **anthropic2025claude**; **google2025gemini**). These developments have intensified concerns about AI's potential to fundamentally alter economic relationships, with prominent researchers warning of catastrophic risks from advanced AI systems (**bengio2024managing**; **bostrom2014superintelligence**; **russell2019human**).

The prospect of a technological singularity—a hypothetical point where AI capabilities undergo explosive, self-reinforcing growth—has moved from science fiction to serious academic discourse (**vinge1993coming**; **good1965speculations**; **kurzweil2024singularity**). While the timing and likelihood of such an event remain deeply uncertain, the possibility raises fundamental questions for asset pricing. How should investors value firms whose capabilities might explode during crises? What happens to standard valuation frameworks when technology exhibits extreme disaster-contingent growth?

This paper develops a consumption-based asset pricing model to analyze these questions. We embed AI stocks in a disaster risk framework where AI capabilities advance during economic crises, reflecting automation demand and labor substitution effects documented in recent empirical work (**zhang2019labor**; **knesl2023automation**; **babina2024artificial**). Our key insight is that AI stocks provide valuable hedging against disaster risk from the household's perspective, but this hedging property becomes destabilizing when sufficiently extreme.

The model yields two main results. First, AI stocks command high valuations due to their negative covariance with consumption growth—they pay off precisely when marginal utility is highest. This formalizes the intuition that investors should be willing to pay premium prices for assets that hedge catastrophic outcomes. Second, beyond a critical threshold of disaster-contingent growth, the model produces infinite valuations, indicating that standard equilibrium cannot be sustained. This mathematical singularity in asset pricing provides a formal representation of concerns that transformative AI could fundamentally destabilize economic relationships.

Our work connects several literatures. The disaster risk literature (**rietz1988equity**; **barro2006rare**; **gabaix2012variable**; **gourio2012disaster**) shows how rare catastrophic events affect asset prices. The labor hedging literature (**santos2006labor**; **benzoni2007portfolio**; **betermier2012hedging**) analyzes how labor income risk influences portfolio choice. Re-

cent work examines AI's impact on firm growth and systematic risk (**babina2024artificial**; **babina2023artificial**). We contribute by explicitly modeling AI's disaster-contingent growth and deriving conditions under which this growth destabilizes asset pricing equilibria.

The remainder of the paper proceeds as follows. Section 2 develops the theoretical model. Section 3 derives the price-dividend ratio and presents quantitative results. Section 4 concludes.

2 Model

We develop a consumption-based asset pricing model with disaster risk to analyze the relationship between AI stocks and singularity concerns.

2.1 Environment

Consider an endowment economy with a representative household having CRRA utility $U(C) = \frac{C^{1-\gamma}}{1-\gamma}$ and time discount factor β. The household's coefficient of relative risk aversion is $\gamma > 0$. Consumption growth follows a simple two-state process:

$$\log(C_{t+1}/C_t) = \begin{cases} 0 & \text{with probability } 1-p \\ -b & \text{with probability } p \end{cases} \tag{1}$$

where $b > 0$ represents disaster severity and $p \in (0, 1)$ is the disaster probability. In the normal state, consumption remains constant. In the disaster state, consumption falls by factor e^{-b}, representing economic crises such as financial panics, pandemics, or major technological disruptions.

This specification follows the disaster risk literature (**barro2006rare**; **gabaix2012variable**), which shows that rare but severe consumption drops can explain key asset pricing puzzles. The two-state structure provides analytical tractability while capturing the essential feature that disasters represent tail events with outsized economic impact.

2.2 AI Asset

The AI sector produces dividends according to:

$$D_t = ae^{hN_t}C_t \tag{2}$$

where $N_t = \sum_{s=1}^{t} \mathbb{1}_{\{\text{disaster at } s\}}$ counts cumulative disasters, $a > 0$ is a scale parameter, and $h > 0$ captures AI's disaster-linked growth. The parameter h is the key innovation in our model—it measures how much AI capabilities advance during each disaster.

This specification reflects several empirical regularities. First, automation demand increases during crises as firms seek to reduce labor costs and improve resilience (**acemoglu2020robots**). Second, AI capabilities may advance more rapidly during disruptions that force organizational change and technology adoption (**babina2024artificial**). Third, labor substitution by capital accelerates during downturns (**zhang2019labor**; **knesl2023automation**).

The exponential form e^{hN_t} captures the possibility of compounding advances—each disaster permanently raises AI's productivity level, and these gains accumulate over time. When $h = 0$, AI dividends simply scale with aggregate consumption, behaving like a standard equity claim. When $h > 0$, AI exhibits disaster-contingent growth, advancing precisely when the economy contracts.

2.3 Pricing Kernel

The stochastic discount factor is:

$$M_{t+1} = \beta \left(\frac{C_{t+1}}{C_t} \right)^{-\gamma} \tag{3}$$

This standard consumption-based pricing kernel (**cochrane2008two**) values assets according to their covariance with consumption growth. Assets that pay off when consumption is low (and marginal utility is high) command premium prices.

The risk-free rate satisfies:

$$R_f = \frac{1}{\mathbb{E}[M_{t+1}]} = \frac{1}{\beta[(1-p) + pe^{\gamma b}]} \tag{4}$$

The risk-free rate decreases in disaster probability p and severity b, reflecting precautionary savings demand. When disasters become more likely or severe, households demand higher returns to hold risky assets, driving down the risk-free rate.

2.4 Key Result: Negative Beta

The AI asset exhibits negative beta with respect to consumption growth. When disasters occur, consumption falls by e^{-b} but AI dividends rise by factor e^h. The covariance is:

$$\text{Cov}\left(\frac{C_{t+1}}{C_t}, \frac{D_{t+1}}{D_t} \right) = p(1-p)(e^{-b} - 1)(e^h - 1) < 0 \tag{5}$$

This negative covariance implies AI stocks hedge disaster risk from the household's perspective. They provide insurance against the worst states of the world, paying off precisely when consumption is lowest and marginal utility is highest. This hedging property drives the valuation results in the next section.

3 AI Stock Valuation

3.1 Deriving the Price-Dividend Ratio

To determine the equilibrium valuation of AI stocks, we solve the Euler equation at time $t = 0$. Let P_t denote the AI stock price. The fundamental asset pricing equation gives:

$$P_t = \mathbb{E}_t[M_{t+1}(D_{t+1} + P_{t+1})] \tag{6}$$

Conjecturing a constant price-dividend ratio $P_t = \phi D_t$ for some $\phi > 0$, we substitute into the asset pricing equation:

$$\phi D_0 = \mathbb{E}_0\left[\beta\left(\frac{C_1}{C_0}\right)^{-\gamma}(D_1 + \phi D_1)\right] = \beta(1+\phi)\mathbb{E}_0\left[\left(\frac{C_1}{C_0}\right)^{-\gamma} D_1\right] \tag{7}$$

The expectation evaluates across two states. In the no-disaster state (probability $1-p$), consumption remains constant at $C_1 = C_0$ and the disaster counter stays at $N_1 = 0$. The dividend is $D_1 = aC_0 = D_0$, and the pricing kernel term is simply 1, yielding $\left(\frac{C_1}{C_0}\right)^{-\gamma} D_1 = D_0$.

In the disaster state (probability p), consumption falls to $C_1 = e^{-b}C_0$ while the disaster counter increments to $N_1 = 1$. The dividend becomes $D_1 = ae^h C_1 = ae^h e^{-b} C_0$. The pricing kernel evaluates to $e^{\gamma b}$, so the full term becomes:

$$\left(\frac{C_1}{C_0}\right)^{-\gamma} D_1 = e^{\gamma b} \cdot ae^h e^{-b} C_0 = e^{\gamma b + h} D_0 \tag{8}$$

Combining these states, the expectation becomes:

$$\mathbb{E}_0\left[\left(\frac{C_1}{C_0}\right)^{-\gamma} D_1\right] = (1-p)D_0 + pe^{\gamma b+h}D_0 = D_0[(1-p) + pe^{\gamma b+h}] \tag{9}$$

Substituting back and canceling D_0 yields:

$$\phi = \beta(1+\phi)[(1-p) + pe^{\gamma b+h}] \tag{10}$$

Defining $\lambda = \beta[(1-p) + pe^{\gamma b+h}]$ and rearranging gives $\phi(1-\lambda) = \lambda$, which solves to:

$$\boxed{\frac{P_0}{D_0} = \phi = \frac{\beta[(1-p) + pe^{\gamma b+h}]}{1 - \beta[(1-p) + pe^{\gamma b+h}]}} \tag{11}$$

This formula is valid only when the denominator remains positive, requiring $\beta[(1-p) + pe^{\gamma b+h}] < 1$. When this condition fails, the price-dividend ratio becomes infinite, indicating that no finite equilibrium price exists.

3.2 Economic Interpretation

The price-dividend ratio depends on three key forces. First, higher disaster severity b increases valuations through two channels: it raises the marginal utility of consumption during disasters (the $e^{\gamma b}$ term) and it increases the value of AI's hedging property. Second, higher AI disaster-linked growth h directly increases valuations by making AI more valuable during crises. Third, these effects interact multiplicatively through the $e^{\gamma b+h}$ term, creating potential for explosive valuations.

The critical threshold occurs when:

$$h \geq -\gamma b + \log\left(\frac{1-\beta(1-p)}{\beta p}\right) \tag{12}$$

Beyond this point, AI's disaster-contingent growth is so extreme that standard valuation breaks down. This mathematical singularity captures the essence of singularity risk in asset markets—when AI capabilities advance too rapidly during crises, the economy cannot sustain equilibrium.

3.3 Quantitative Results

Table 1 presents price-dividend ratios for various combinations of disaster severity b and AI disaster-linked growth h, using baseline parameters $\gamma = 2$, $\beta = 0.96$, and $p = 0.01$. These parameters are standard in the disaster risk literature (**barro2006rare**).

The table reveals several striking patterns. When AI exhibits no disaster-linked growth ($h = 0$), the price-dividend ratio remains constant at 23.76 regardless of disaster severity, reflecting standard consumption-based pricing without hedging benefits. This baseline corresponds to a standard equity claim that simply scales with aggregate consumption.

As h increases, valuations rise substantially, with the effect amplifying at higher disaster severities. For moderate AI growth parameters ($h \leq 1.0$), price-dividend ratios range from approximately 24 to 291, representing elevated but finite valuations. These valuations reflect

Table 1: Price-Dividend Ratios Under Disaster Risk

$b \setminus h$	0.0	0.5	1.0	1.5	2.0	2.5	3.0
0.1	23.76	27.38	32.13	38.56	47.63	61.29	83.88
0.2	23.76	30.15	39.94	56.35	88.46	162.86	442.50
0.3	23.76	33.57	51.26	92.45	213.75	1071.00	∞
0.4	23.76	37.86	69.23	165.00	765.00	∞	∞
0.5	23.76	43.29	100.71	376.67	∞	∞	∞
0.6	23.76	50.25	160.71	1620.00	∞	∞	∞
0.7	23.76	59.40	291.00	∞	∞	∞	∞

Note: This table shows price-dividend ratios ϕ for various combinations of disaster severity b and AI disaster-linked growth h. Parameters: $\gamma = 2$, $\beta = 0.96$, $p = 0.01$. The symbol ∞ indicates parameter combinations where $\beta[(1-p) + pe^{\gamma b+h}] \geq 1$, implying no finite equilibrium exists.

AI's hedging value—investors are willing to pay premium prices for assets that pay off during disasters.

The most dramatic feature emerges at higher values of h. The model produces infinite valuations when $\beta[(1-p) + pe^{\gamma b+h}] \geq 1$, which occurs approximately when $h \geq 4.02 - 2b$. These infinite entries represent parameter combinations where the economy cannot sustain equilibrium. When AI's disaster-contingent growth becomes sufficiently extreme, the standard asset pricing framework breaks down entirely.

This breakdown captures the essence of singularity risk in financial markets. The critical threshold decreases with disaster severity: for $b = 0.1$, the model remains stable until $h \approx 3.82$, while for $b = 0.5$, instability emerges at $h \approx 3.02$. Economically, this reflects that more severe disasters require less extreme AI growth to trigger valuation explosions, since the interaction term $e^{\gamma b+h}$ compounds both effects.

The negative covariance between AI dividends and consumption growth creates a powerful hedging demand. During disasters, when consumption falls by factor e^{-b} and marginal utility spikes by $e^{\gamma b}$, AI dividends rise by e^{h}. Households value this insurance property, bidding up AI stock prices. However, when this hedging benefit becomes too strong relative to the discount rate, the transversality condition fails and no finite equilibrium price exists.

3.4 Implications for Asset Markets

Our results suggest that current AI stock valuations may reflect rational pricing of disaster hedging benefits rather than speculative bubbles. If investors believe AI capabilities will advance during crises—through automation, labor substitution, or technological breakthroughs forced by disruption—then high price-dividend ratios are justified by standard asset pricing

theory.

However, the model also reveals fundamental limits to this logic. Beyond a critical threshold of disaster-contingent growth, standard valuation frameworks break down. This mathematical singularity provides a formal representation of concerns that transformative AI could fundamentally destabilize economic relationships (**vinge1993coming**; **bostrom2014superintelligence**).

The model's predictions could be tested empirically by examining how AI stock returns covary with consumption growth during crises. If AI stocks exhibit negative beta during disasters—rising when consumption falls—this would support the hedging interpretation. Conversely, if AI stocks crash alongside consumption during crises, this would suggest current valuations reflect speculation rather than rational hedging demand.

4 Conclusion

We have developed a consumption-based asset pricing model where AI capabilities advance during economic disasters. The model yields two main insights. First, AI stocks provide valuable hedging against disaster risk, commanding high valuations due to their negative covariance with consumption growth. This formalizes the intuition that investors should be willing to pay premium prices for assets that pay off during catastrophic outcomes.

Second, when AI's disaster-contingent growth becomes sufficiently extreme, standard asset pricing breaks down entirely. The model produces infinite valuations beyond a critical threshold, indicating that no finite equilibrium can be sustained. This mathematical singularity in asset pricing provides a formal representation of concerns about transformative AI destabilizing economic relationships.

Our framework connects recent advances in AI capabilities (**chollet2024arc**; **deepseek2025deepseek**) with longstanding concerns about technological singularities (**vinge1993coming**; **good1965speculations**). While the timing and likelihood of transformative AI remain deeply uncertain, our model shows how such possibilities affect asset prices today. Investors pricing AI stocks must balance the hedging value of disaster-contingent growth against the risk that such growth becomes destabilizing.

Several extensions merit further research. First, incorporating endogenous AI development decisions could illuminate how investment incentives change near the singularity threshold. Second, analyzing portfolio choice in this environment could reveal how households optimally hedge singularity risk. Third, extending the model to multiple AI firms with heterogeneous disaster-contingent growth could explain cross-sectional variation in AI stock valuations.

The rapid advancement of AI capabilities makes these questions increasingly urgent. As

reasoning models approach human-level performance on complex tasks (**anthropic2025claude**; **google2025gemini**), and concerns about catastrophic AI risks intensify (**bengio2024managing**; **jones2024ai**), understanding how financial markets price transformative technology becomes essential. Our model provides a first step toward formalizing these concerns within standard asset pricing theory.

A A Purely Human Perspective

The following is the README.md file from the GitHub repository:

```
# Prompts-to-Paper

Writes a paper about hedging a negative AI singularity, using AI.

- make-paper.py writes a paper

- plan0408-piecewise.yaml contains the prompts

- make-many-papers.py runs make-paper.py many times.

The README is entirely human-written.  Please forgive typos and errors.

-Andrew Chen, April 9, 2025

# Motivation

On March 8, 2025 I thought I should write a paper about hedging the AI
singularity.

I was worked up.  I had been repeatedly shocked by AI progress.  I was using
AI to prove theorems, vibe coding, and AI lit reviews in my daily life.  Six
months ago, I had thought each of these things is impossible.

What will happen in the next six years?!  Will my entire job be replaced by
AI? I have no idea.

But I do know that if there are huge AI disruptions, then tech stocks will
most likely benefit.  So if anything bad happens to my human capital, I
could at least partially hedge.  Strangely, I hadn't heard about this concept
before.

I asked a friend if he would be interested in working on this paper.
Unfortunately, he was busy.  So, I thought I should use AI to write the
paper.

## Inspiration

This project was inspired by Novy-Marx and Velikov (2025) and Chris Lu et
al.  (2024).  These projects use AI to generate massive amounts of academic
research.  My goal differs in quality over quantity.  I want to generate just
one paper, but one paper that (I hope) people find is worth reading.
```

This project was also inspired by Garleanu, Kogan, and Panageas's (2012) beautiful model of innovation and displacement risk. I read Garleanu et al. back when I was a PhD student and it has stuck with me.

Last, I drew from Hadfield-Menell and Hadfield (2018) and Bengio (2023), who apply ideas from economics to AI catastrophe risk. Hadfield-Menell and Hadfield (2018) explains the connection between incomplete contracting and AI alignment. Bengio (2023) frames AI catastrophe risk in terms of what I would call decision theory and human incentives---though the essay is written in plain English.

Previously, I dismissed "AI safety" as a politicized, activist cause. I still don't like the term "AI safety." "Safety" is such an excuse for exercising power over others.

But then the nature of AI changed. Things were progressing much faster than I expected. The Jan 15, 2025 episode of Machine Learning Street Talk with Yoshua Bengio left an impression on me. Bengio talked about AI catastrophe risk with no activism, no fear mongering. It was a straight, rational discussion of the seriousness of AI catastrophe risk.

The Paper Generation Process

Is this paper really written by AI?

I'd say the AI are junior co-authors.

If they were human, I would absolutely have to give o1, Sonnet, and ChatGPT Deep Research credit as co-authors. They did the math, writing, and literature reviews. Sonnet also wrote most of the code (via the Cursor AI IDE).

Of course, the prompts show I did substantial hand-holding. The Github commits show even more human labor. They tell the story of me getting to know my "co-authors." It was hard to communicate style issues and LaTeX instructions, leading to many, many commits.

To be honest, writing this paper would have been much easier if I had done more of the work myself.

But that can happen with human co-authors too.

Perhaps in the next few years, AI and humans will be equal co-authors.

Paper Iterations

Like human-written papers, the writing process was iterative. The first formalizations were terrible.

`plan0313-laborshare.yaml` (from March 13) contains prompts for a neoclassical growth model, where the capital share suddenly increases. ChatGPT-o1 patiently explained to me why this is a bad model.

me: > I thought there would be a wage risk effect that leads to higher investment for the more risk averse agent. High capital share means low or even no wage income.

ChatGPT-o1: > Below is an explanation of why one might **expect** a "wage-risk" channel in which **more** risk aversion could lead to **more** *additional* investment (relative to the no-jump benchmark) in the event that the capital share might jump to 1. However, this channel **does not operate** in the usual **representative-agent** version of the model---there, wage and capital income ultimately go to the *same* agent, so there is no meaningful "hedge" of wage risk.

I went through several iterations of the model with Claude 3.7 Sonnet (thinking mode) and ChatGPT-o1. The only derivations I did myself were to check o1's work, which I found to be quite reliable.

The final `plan0408-piecewise.yaml` uses a simplified Barro-Rietz disaster model, with two agents (though only one is relevant for stock prices). I slowly walk the AIs through the writing, using ten prompts, to maintain the writing quality.

Literature Reviews

A key step was generating lit reviews (`./lit-context/`) which were used as context in the prompts. I made lit reviews using ChatGPT's Deep Research (launched Feb 2025) until I ran out of credits. I used Claude Web Search (launched March 20, 2025) for the remainder.

These new products are a game changer. Both Novy-Marx and Velikov (2025) and Chris Lu et al. (2024) ran into hallucinated citations. Building off OpenAI Deep Research and Claude Web Search led to zero hallucinations in the paper drafts.

Still, I would occassionaly run into mis-citations. Every 15 to 20

citations, I would see a cite that mis- or overinterprets the cited paper. I suppose that's not so different than the human-written citation error rate. But I hate finding misinterpretations in the literature so I purposefully limited the number of cites in the paper.

AI Model Selection

I found it best to use o1 for theory, and Sonnet for writing. It's well-known that these are the strengths of these two models.

I briefly tried having Llama 3.1 405b do the writing. It was terrible! The cutting edge models seemed necessary.

I did not try many other models, in order to get this paper out quickly. Gemini 2.5's release, at the end of March 2025, was *hype*. I tried it out briefly and was impressed. But I gritted my teeth and ignored it. AI progress is too fast to update this paper continuously.

Picking the best of N papers

The writing quality varies across each run. Rather than try to prompt engineer an error free, insightful paper, I decided to just generate N papers and choose the best one.

5 drafts of the paper can be found in `./manyout0408-pdf/`. They're broadly similar and I would be willing to put my name on most of them.

I ended up choosing `paper-run-05.pdf`. The abstract overemphasizes the model, but otherwise I like the paper. I hope you find the paper worth reading.

Lessons about Research

A common response to Novy-Marx and Velikov (2025) is: "people are not ready for this." I heard concerns that peer review will be inundated with AI-generated slop.

Working on this paper gave me a different perspective. It made me think about the fundamentals. I think the fundamentals are the following:

1. Readers want to learn something interesting and true.

2. Readers don't want to check all the math.

3. A system of author reputations makes 1 and 2 possible.

AI-generated papers don't change any of these fundamentals. Critically, fundamental 3 made me quite wary of putting my name on AI slop. As a result, I don't think AI-generated papers will change much about peer review, at least not the current generation of AI.

Limitations of the Current AI (April 9, 2025)

This will likely be out of date by the time you read it.

But right now, AI is like a junior co-author with a talent for mathematics and elegant writing, but sub-par economics reasoning.

For example, Sonnet often fails to recognize that an economic model does not capture an important channel. This is common (no model can capture everything). The standard practice is to dance gingerly around the channel in the writing. A decent PhD student can recognize this. But Sonnet cannot.

AI also struggles to generate a satisfying economic model on its own. Its models are typically either too simplistic or too complicated.

I opted not to add empirical work or numerically-solved models. AI can do both, but both require connecting to the outside world, and a plethora of technical challenges.

Last, I often found the literature discussions insufficiently careful (as discussed above).

But how long will these limitations last?

The Future of AI and Economics Research (Speculative)

At some point, 2024-style economic analysis will be "on tap." You'll be able to go to a chatbot and ask "write me a paper about hedging AI disaster risk," and it will return you something like this paper (probably something much better).

"Economics on tap" could be a disaster for the economics labor market (could be). It certainly *will* be an extremely cheap substitute for at least some economists' labor. I suppose the questions is whether that will result in a strong substitution away from labor.

The optimistic argument is that AI also *complements* economists' labor. Perhaps, the number of economists will remain the same, but our research output increases in terms of both quantity and quality.

```
But there are reasons why total research output is limited.  Two key factors
in academic publishing are attention and reputation (Klamer and van Dalen
2001, J of Economic Methodology).  Readers can only pay attention to so
many scholars.  These scholars, in turn, can only pay attention to so may
projects.

I'm not saying that I *expect* a disaster for the economics labor market.  But
even if it's highly unlikely, it's still a scenario that economists should consider.
```

B Prompts Used to Generate This Paper

Each prompt consists of context and instructions. The context consists of the responses to the previous prompts, and may include literature reviews (all AI generated). For writing tasks (using Claude 3.7 Sonnet), a system prompt is also included.

For further details, see `https://github.com/chenandrewy/Prompts-to-Paper/`.

The system prompt and instructions are listed below.

System Prompt (model: anthropic.claude-sonnet-4-5-20250929-v1:0)

```
You are writing an economics paper about AI stocks and risk. Write
   clearly and use math when needed. Output latex format.
```

Instruction: 01-model-prose (model: anthropic.claude-sonnet-4-5-20250929-v1:0)

```
Describe a model with two agents (AI owners and households),
   consumption growth with disaster risk, and an AI asset. The model
    should show how AI stocks relate to singularity risk.

Assumptions:
- Representative household with CRRA = \\gamma, time preference = \\
   beta
- Log consumption growth = 0 normally, = -b in disaster (probability
    p)
- AI asset dividend: D\\_t = a e\\^\\{h N\\_t\\} C\\_t where N\\_t
   counts disasters
- Disasters are bad for households but good for AI stocks

Keep it brief. No sectioning commands.
```

Instruction: 02-result-notes (model: anthropic.claude-sonnet-4-5-20250929-v1:0)

```
Derive the price/dividend ratio for the AI asset at t=0. Show your
   work but be concise.
```

Instruction: 03-table-notes (model: anthropic.claude-sonnet-4-5-20250929-v1:0)

```
Make a table showing price/dividend ratios for different parameter
   values. Use gamma=2, beta=0.96, p=0.01. Try different values of b
    and h.
```

Instruction: 04-resultandtable-prose (model: anthropic.claude-sonnet-4-5-20250929-v1:0)

```
Convert the math from `02-result-notes` and `03-table-notes` into
   prose. This should follow the model description. Show the key
   formula and table. Explain what they mean for AI stock valuations
   .

Keep it concise. No bullet points.
```

Instruction: 10-full-paper (model: anthropic.claude-sonnet-4-5-20250929-v1:0)

```
Write a complete paper titled "Hedging the AI Singularity" using all
    the `*-prose` context.

The paper should have:
- Abstract (50 words): Main point is that AI stocks hedge against
   negative AI singularity
- Introduction (brief): AI is advancing fast, this affects stock
   prices
- Model section: Use the prose from previous prompts
- Conclusion (brief): Summarize main findings

Output complete latex document with preamble. Use `template.tex` as
   template.
```